



Software

# Introduction to Convolutional Neural Networks

Modified from Intel's original slides: ' Introduction to Convolutional Neural Networks '

# Legal Notices and Disclaimers

This presentation is for informational purposes only. INTEL MAKES NO WARRANTIES, EXPRESS OR IMPLIED, IN THIS SUMMARY.

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. Check with your system manufacturer or retailer or learn more at [intel.com](https://www.intel.com).

This sample source code is released under the [Intel Sample Source Code License Agreement](#).

Intel and the Intel logo are trademarks of Intel Corporation in the U.S. and/or other countries.

\*Other names and brands may be claimed as the property of others.

Copyright © 2017, Intel Corporation. All rights reserved.

# Motivation – Image Data

- Traditional neural networks treat all inputs as interchangeable, ignoring relationships between individual inputs.
- Inputs are processed as an ordered set of variables, lacking spatial context.
- Goal: Integrate domain knowledge into the neural network architecture, to better capture structure and relationships in image data.

# Motivation



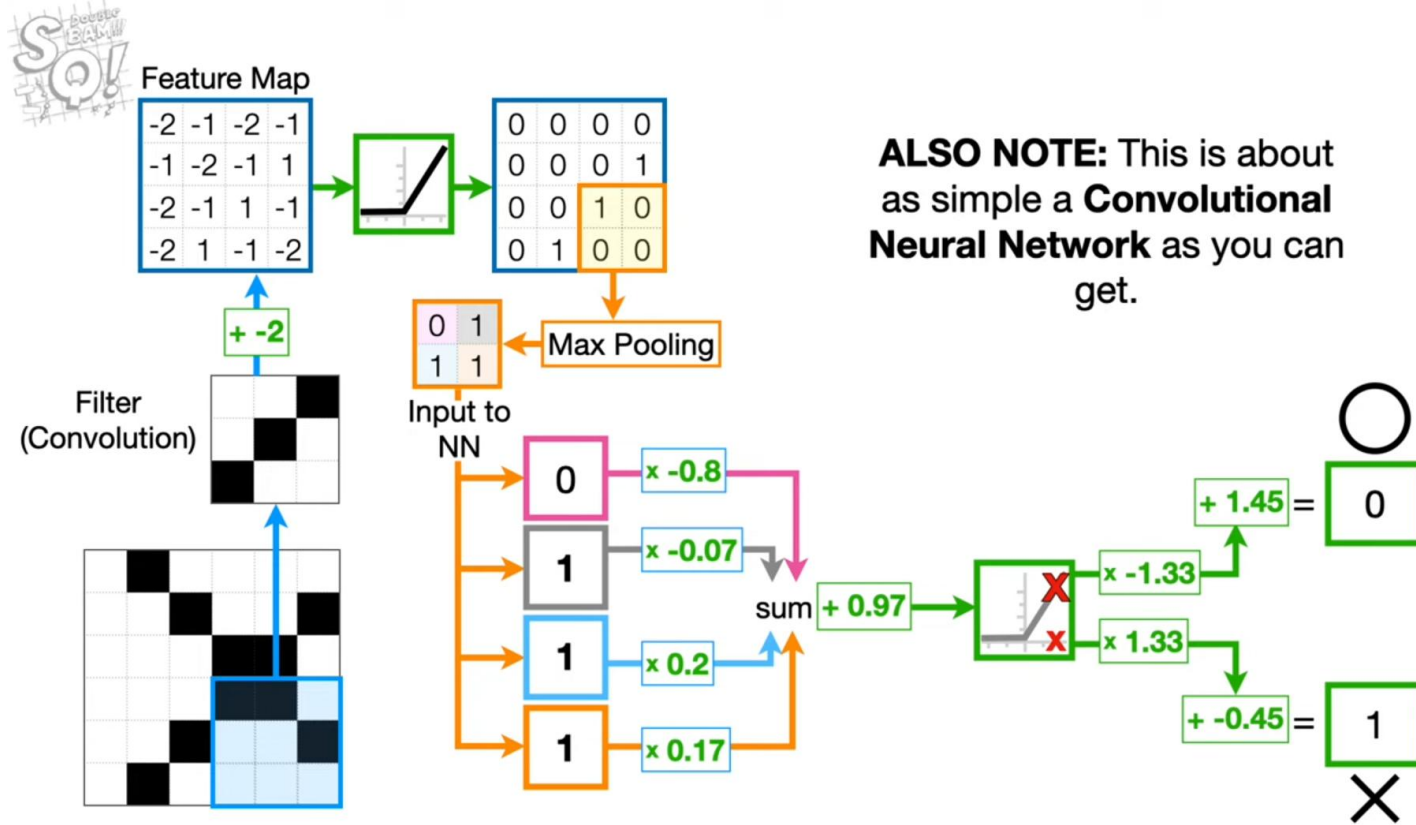
- Image data has important structures, such as:
  - **“Topology”** of pixels: spatial relationships between pixels
  - **Local Similarity**: nearby pixels often have similar values
  - **Edges and Shapes**: fundamental building blocks of images
  - **Translation Invariance**: objects retain meaning when shifted
  - **Scale Invariance**: objects may appear at different sizes
  - **Lighting and Contrast**: variations affect pixel intensity
  - **Human Visual System Insights**: leverage natural perception

# Motivation – Image Data

- Fully connected networks require an impractical number of parameters for image data.
- MNIST images ( $28 \times 28$  grayscale) are manageable, but typical images (e.g.,  $200 \times 200$  RGB) have  $\sim 120,000$  input values.
- A single fully connected layer would require  $(200 \times 200 \times 3)^2 = 14,400,000,000$  weights! ( $\sim 14.4$  billion weights).
- High parameter count increases variance (overfitting).
- Convolutional layers reduce parameters by focusing on local patterns, introducing "bias" into the network.

# Image Classification with Convolutional Neural Networks (CNNs)

(from StatQuest with Josh Starmer)



# Motivation

- Features in images are hierarchically composed.
  - Edges  $\rightarrow$  Shapes  $\rightarrow$  Relationships between shapes  $\rightarrow$  Textures.
- Example: Recognizing a cat.
  - **Cat**: two eyes in a specific arrangement + fur texture.
  - **Eyes**: dark circle (pupil) inside another circle.
  - **Circle**: combination of edge detectors.
  - **Fur**: repeating edge patterns.

# Kernels (Filters)

- A *kernel* is a grid of weights applied to an image, centered on a pixel.
- Each weight is multiplied by the corresponding pixel value beneath it.
- Output for the centered pixel:

$$\sum_{p=1}^P W_p \cdot pixel_p$$

- Kernels are fundamental in traditional image processing techniques:
  - Blur
  - Sharpen
  - Edge detection
  - Emboss



# Kernel: 3x3 Example

Input

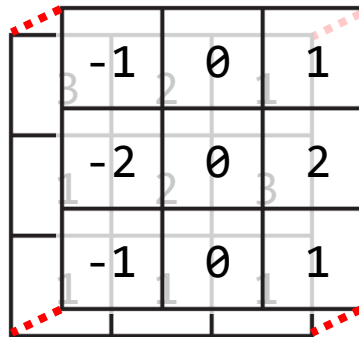
3	2	1
1	2	3
1	1	1

Kernel

-1	0	1
-2	0	2
-1	0	1

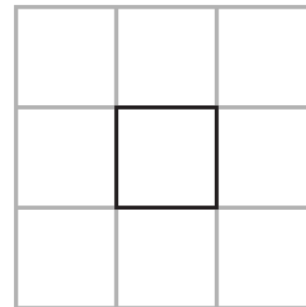
Output


# Kernel: 3x3 Example



	-1	0	1	
	-2	0	2	
	-1	0	1	

Output




# Kernel: 3x3 Example

Input			Kernel			Output		
3	2	1	-1	0	1			
1	2	3	-2	0	2		2	
1	1	1	-1	0	1			

$$\begin{aligned} &= (3 \cdot -1) + (2 \cdot 0) + (1 \cdot 1) \\ &+ (1 \cdot -2) + (2 \cdot 0) + (3 \cdot 2) \\ &+ (1 \cdot -1) + (1 \cdot 0) + (1 \cdot 1) \end{aligned}$$

$$= -3 + 1 - 2 + 6 - 1 + 1 = 2$$

# Kernels as Feature Detectors

Can think of kernels as a "local feature detectors"

Vertical Line Detector

-1	1	-1
-1	1	-1
-1	1	-1

Horizontal Line Detector

-1	-1	-1
1	1	1
-1	-1	-1

Corner Detector

-1	-1	-1
-1	1	1
-1	1	1

Edge Detection

0	1	0
1	-4	1
0	1	0

"Strong" Edge Detection

-1	-2	-1
0	0	0
1	2	1

Try Convolution Yourself

visit <https://setosa.io/ev/image-kernels>


# Convolutional Neural Nets

Primary Ideas behind Convolutional Neural Networks:


- Neural networks learn optimal kernels during training.
- Apply the same kernels across the entire image (translation invariance).
- Reduces parameters and minimizes overfitting by controlling variance (bias-variance tradeoff).

# Convolutions

An Input RGB Image



4 Feature Maps  
(from 4 Filters)




# Convolution Settings – Grid Size or Kernel Size

## Grid Size or Kernel Size (Height x Width):

- Number of pixels the kernel processes at once.
- Typically use odd numbers to ensure a “center” pixel.
- Kernels can be rectangular, not necessarily square.

Height: 3, Width: 3



Height: 1, Width: 3



Height: 3, Width: 1



# Convolution Settings - Padding

## Padding

- Kernels cause an “edge effect” where edge pixels are not used as “center pixels” since there are not enough surrounding pixels
  - "center pixels" refer to the pixels in the input over which the kernel is centered
- Padding adds extra pixels around the image to include all original pixels as center pixels.
- Commonly used padding: zero-padding (added pixels have a value of zero), replication-padding (added pixels replicate the values of the nearest edge pixels).



# Without Padding

padding="valid" --> no padding

padding="same" --> output size is the same as of input

1	2	0	3	1
1	0	0	2	2
2	1	2	1	1
0	0	1	0	0
1	2	1	1	1

input

-1	1	2
1	1	0
-1	-2	0

kernel

-2		

output

$$\text{output size} = n - f + 1$$

# With Padding

padding="valid" --> no padding

padding="same" --> output size is the same as of input

0	0	0	0	0	0	0
0	1	2	0	3	1	0
0	1	0	0	2	2	0
0	2	1	2	1	1	0
0	0	0	1	0	0	0
0	1	2	1	1	1	0
0	0	0	0	0	0	0

original input with zero-padding

-1	1	2
1	1	0
-1	-2	0

kernel

-1				

output

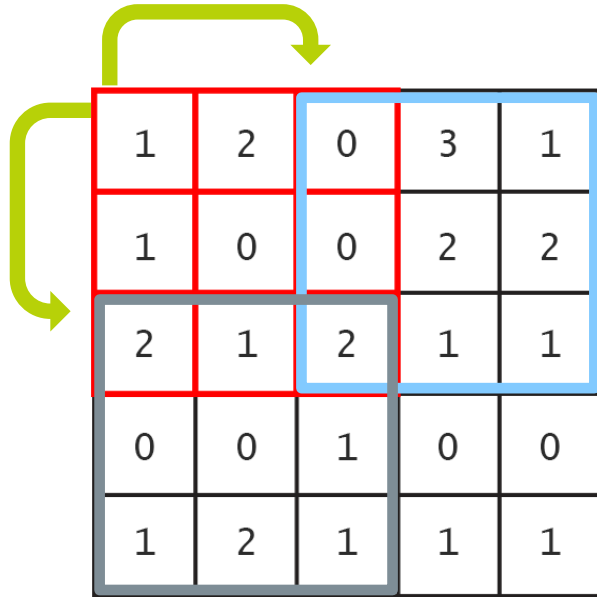
$$\text{output size} = n + 2p - f + 1$$

# Convolution Settings

## Stride

- The step size for moving the kernel across the image.
- Can differ for vertical and horizontal steps (but typically the same).
- Larger strides ( $>1$ ) reduce the output dimensions, effectively downsampling the **feature map**.
  - A **feature map** is the output of a convolutional layer in a neural network.
  - It represents the spatial activation of features detected by a specific filter (kernel) applied to the input data.

# Stride 2 Example – No Padding



1	2	0	3	1
1	0	0	2	2
2	1	2	1	1
0	0	1	0	0
1	2	1	1	1

input

-1	1	2
1	1	0
-1	-2	0


kernel

-2	

output

$$\text{output size} = \left\lfloor \frac{n - f}{s} + 1 \right\rfloor$$

# Stride 2 Example – With Padding



-1	1	2
1	1	0
-1	-2	0

kernel

-1	2	
3		

output

$$\text{output size} = \left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor$$

# Convolutional Settings - Depth

- Images can have multiple values per pixel, called “channels” (e.g., RGB = 3 channels, CMYK = 4 channels).
- The number of channels is the depth of the image.
- The kernel itself will also have a “depth”, the same size as the number of input channels
  - A kernel's depth matches the input's depth to process all channels simultaneously.
- The kernel performs a dot product between two 3D grids (input and kernel).
- Example: a 5x5 kernel on an RGB image has  $5 \times 5 \times 3 = 75$  weights.

# Convolutional Settings - Depth

- The output depth equals the number of kernels applied in the layer.
- Each kernel produces one feature map.
- Multiple kernels allow the network to learn diverse features, such as edges, textures, and patterns, at different orientations and scales.
- Example: If a layer uses 10 kernels, the output depth will be 10.


# Pooling

- **Purpose:** reduces image size by mapping a patch of pixels to a single value.
- **Effect:** shrinks the spatial dimensions of the feature map.
- **Key Feature:** no learnable parameters; applies predefined operations.
- **Types:** common pooling operations include max pooling and average pooling.




# Pooling: Max-Pool

- For each distinct patch, represent it by the maximum
- 2x2 maxpool shown below



# Pooling: Average-Pool

- For each distinct patch, represent it by the average
- 2x2 avgpool shown below.





Software